

## Retail Data Analysis

### Code Logic

#### Logic for Python Script 'spark-streaming.py'

Setting up the system dependencies for Cloudera distribution by importing necessary libraries, modules and the path variables

```
import os
import sys

os.environ["PYSPARK_PYTHON"] = "/opt/cloudera/parcels/Anaconda/bin/python"
os.environ["JAVA_HOME"] = "/usr/java/jdk1.8.0_232-cloudera/jre"
os.environ["SPARK_HOME"]="/opt/cloudera/parcels/SPARK2-2.3.0.cloudera2-1.cdh5.13.3.p0.316101/lib/spark2/"
os.environ["PYLIB"] = os.environ["SPARK_HOME"] + "/python/lib"
sys.path.insert(0, os.environ["PYLIB"] + "/py4j-0.10.6-src.zip")
sys.path.insert(0, os.environ["PYLIB"] + "/pyspark.zip")

from pyspark.sql import SparkSession
from pyspark.sql.functions import *
from pyspark.sql.types import *
```

#### Writing the Python functions, which contain the logic for the UDFs

1. Total Cost UDF - To calculate the total income from every invoice I needed to calculate the income from sale of each product, so I multiplied the unit price of the product with the quantity of the product purchased. The sum of this cost across the products in that invoice gives me the total cost of the order. I also made sure that if the transaction is a return transaction, the total cost is negative.

```
def find_total_order_cost(items, trn_type):
    if items is not None:
        total_cost = 0
        item_price = 0
        for item in items:
            item_price = (item['quantity'] * item['unit_price'])
            total_cost = total_cost + item_price
            item_price = 0

        if trn_type == "RETURN":
            return total_cost * -1
        else:
            return total_cost
```

2. Total Items UDF - To calculate the number of products in every invoice I added the quantity ordered of each product in that invoice

```
def find_total_item_count(items):
    if items is not None:
        total_count = 0
        for item in items:
            total_count = total_count + item['quantity']
        return total_count
```

3. Is Order UDF - To determine if invoice is for an order or not I used an if-else statement

```
def flag_isOrder(trn_type):  
    if trn_type == "ORDER":  
        return(1)  
    else:  
        return(0)
```

4. Is Return UDF - To determine if invoice is for a return or not I used an if-else statement

```
def flag_isReturn(trn_type):  
    if trn_type == "RETURN":  
        return(1)  
    else:  
        return(0)
```

**Initialising the Spark session and setting the log level to error as a good practice**

```
spark = SparkSession \  
    .builder \  
    .appName("spark-streaming") \  
    .getOrCreate()  
spark.sparkContext.setLogLevel('ERROR')
```

**Reading input data from Kafka mentioning the details of the Kafka broker, such as bootstrap server, port and topic name**

```
orderRawData = spark.readStream \  
    .format("kafka") \  
    .option("kafka.bootstrap.servers", "18.211.252.152:9092") \  
    .option("startingOffsets", "earliest") \  
    .option("failOnDataLoss", "false") \  
    .option("subscribe", "real-time-project") \  
    .load()
```

**Defining JSON schema of each order, using appropriate datatypes and StructField in the case of the item attributes**

```
jsonSchema = StructType() \  
    .add("invoice_no", LongType()) \  
    .add("country", StringType()) \  
    .add("timestamp", TimestampType()) \  
    .add("type", StringType()) \  
    .add("items", ArrayType(StructType([  
        StructField("SKU", StringType()),  
        StructField("title", StringType()),  
        StructField("unit_price", FloatType()),  
        StructField("quantity", IntegerType()),  
    ])))
```

**Reading the raw JSON data from Kafka as 'order stream' by casting it to string and storing it into the alias 'data'**

```
orderStream = orderRawData.select(from_json(col("value").cast("string"),  
jsonSchema).alias("data")).select("data.*")
```

**Defining the UDFs by Converting the Python functions I defined earlier, and assigning the appropriate return datatype**

```
sum_total_order_cost = udf(find_total_order_cost, FloatType())
sum_total_item_count = udf(find_total_item_count, IntegerType())
sum_isOrder = udf(flag_isOrder, IntegerType())
sum_isReturn = udf(flag_isReturn, IntegerType())
```

**Calculating the additional columns according to the required input values**

```
expandedOrderStream = orderStream \
    .withColumn("total_cost", sum_total_order_cost(orderStream.items,
orderStream.type)) \
    .withColumn("total_items", sum_total_item_count(orderStream.items)) \
    .withColumn("is_order", sum_isOrder(orderStream.type)) \
    .withColumn("is_return", sum_isReturn(orderStream.type))
```

**Writing the summarised input values to console, using 'append' output method and applying truncate as false and setting the processing time to 1 minute**

```
extendedOrderQuery = expandedOrderStream \
    .select("invoice_no", "country", "timestamp", "total_cost", "total_items",
"is_order", "is_return") \
    .writeStream \
    .outputMode("append") \
    .format("console") \
    .option("truncate", "false") \
    .trigger(processingTime = "1 minute") \
    .start()
```

**Calculating time-based KPIs (Total sale volume, OPM, Rate of return, Average transaction size) having tumbling window of one minute and watermark of one minute.**

```
aggStreamByTime = expandedOrderStream \
    .withWatermark("timestamp", "1 minute") \
    .groupBy(window("timestamp", "1 minute", "1 minute")) \
    .agg(sum("total_cost").alias("total_sale_volume"),
        count("invoice_no").alias("OPM"),
        avg("is_return").alias("rate_of_return"),
        avg("total_cost").alias("average_transaction_size")
    ) \
    .select("window", "OPM", "total_sale_volume", "average_transaction_size",
"rate_of_return")
```

**Writing the time-based KPIs data to HDFS - HDFS into JSON files for each one-minute window, using 'append' output mode, setting truncate as false, and specifying the HDFS output path for both the KPI files and for their checkpoints. Ten 1-minute window batches were taken.**

```
queryByTime = aggStreamByTime.writeStream \
    .format("json") \
    .outputMode("append") \
    .option("truncate", "false") \
    .option("path", "/user/ec2-user/time_kpi") \
    .option("checkpointLocation", "/user/ec2-user/time_kpi_checkpoints") \
    .trigger(processingTime="1 minute") \
    .start()
```

Calculating time-and-country-based KPIs (Total sale volume, OPM, Rate of return) having tumbling window of one minute and watermark of one minute. Here I grouped by window and country both.

```
aggStreamByCountry = expandedOrderStream \
    .withWatermark("timestamp", "1 minute") \
    .groupBy(window("timestamp", "1 minute", "1 minute"), "country") \
    .agg(sum("total_cost").alias("total_sale_volume"),
        count("invoice_no").alias("OPM"),
        avg("is_return").alias("rate_of_return")) \
    .select("window", "country", "OPM", "total_sale_volume", "rate_of_return")
```

Writing the the time-and-country-based KPIs data to HDFS into JSON files for each one-minute window, using 'append' output mode, setting truncate as false, and specifying the HDFS output path for both the KPI files and for their checkpoints. Ten 1-minute window batches were taken.

```
queryByCountry = aggStreamByCountry.writeStream \
    .format("json") \
    .outputMode("append") \
    .option("truncate","false") \
    .option("path", "/user/ec2-user/country_kpi") \
    .option("checkpointLocation", "/user/ec2-user/country_kpi_checkpoints") \
    .trigger(processingTime="1 minute") \
    .start()
```

Indicating Spark to await termination

```
extendedOrderQuery.awaitTermination()
queryByCountry.awaitTermination()
queryByTime.awaitTermination()
```

## Console Commands

I started by logging into the ec2 instance as 'ec2-user'

Next, I downloaded the Spark-SQL-Kafka jar file. This jar is used to run the Spark Streaming-Kafka codes

```
wget https://ds-spark-sql-kafka-jar.s3.amazonaws.com/spark-sql-kafka-0-10_2.11-2.3.0.jar
```

Next, I created the 'spark-streaming.py' file having the code discussed above

```
vi spark-streaming.py
```

Next, I set the Kafka Version using the following command

```
export SPARK_KAFKA_VERSION=0.10
```

Finally, I ran the spark2-submit command, specifying the jar and python file

```
spark2-submit --jars spark-sql-kafka-0-10_2.11-2.3.0.jar spark-streaming.py
```



## Example table - Final Summarised Input Values

Batch: 0

invoice_no	country	timestamp	total_cost	total_items	is_order	is_return
154132546539194	United Kingdom	2021-12-30 18:51:23	70.08	56	1	0
154132546539195	United Kingdom	2021-12-30 18:51:24	33.79	41	1	0
154132546539196	United Kingdom	2021-12-30 18:51:30	-44.88	30	0	1
154132546539197	United Kingdom	2021-12-30 18:51:34	79.47	39	1	0
154132546539198	United Kingdom	2021-12-30 18:51:40	14.22	6	1	0
154132546539199	United Kingdom	2021-12-30 18:51:41	1.25	1	1	0
154132546539200	United Kingdom	2021-12-30 18:51:47	1.65	1	1	0
154132546539201	Netherlands	2021-12-30 18:51:57	59.25	27	1	0
154132546539202	United Kingdom	2021-12-30 18:51:57	68.15	41	1	0
154132546539203	United Kingdom	2021-12-30 18:51:58	5.79	1	1	0
154132546539204	United Kingdom	2021-12-30 18:51:59	2.48	2	1	0
154132546539205	United Kingdom	2021-12-30 18:52:16	28.33	33	1	0
154132546539206	United Kingdom	2021-12-30 18:52:23	32.34	13	1	0
154132546539207	United Kingdom	2021-12-30 18:52:28	7.5	2	1	0
154132546539208	United Kingdom	2021-12-30 18:52:36	31.5	22	1	0
154132546539209	United Kingdom	2021-12-30 18:52:39	17.849998	3	1	0
154132546539210	United Kingdom	2021-12-30 18:53:08	-38.72	7	0	1
154132546539211	Switzerland	2021-12-30 18:53:08	68.0	32	1	0
154132546539212	United Kingdom	2021-12-30 18:53:09	75.06	52	1	0
154132546539213	United Kingdom	2021-12-30 18:53:27	123.38	71	1	0

only showing top 20 rows

I checked HDFS to make sure the KPI files were present

```
hadoop fs -ls /user/ec2-user
```

```
[ec2-user@ip-10-0-0-71 ~]$ hadoop fs -ls /user/ec2-user
Found 6 items
drwx----- - ec2-user ec2-user      0 2021-12-30 11:02 /user/ec2-user/.Trash
drwxr-xr-x - ec2-user ec2-user      0 2021-12-30 11:30 /user/ec2-user/.sparkStaging
drwxr-xr-x - ec2-user ec2-user      0 2021-12-30 11:30 /user/ec2-user/country_kpi
drwxrwxrwx - ec2-user ec2-user      0 2021-12-30 11:11 /user/ec2-user/country_kpi_checkpoints
drwxr-xr-x - ec2-user ec2-user      0 2021-12-30 11:30 /user/ec2-user/time_kpi
drwxrwxrwx - ec2-user ec2-user      0 2021-12-30 11:11 /user/ec2-user/time_kpi_checkpoints
```

I also checked the folders to see the JSON files

```
hadoop fs -ls /user/ec2-user/time_kpi/
```

```
[ec2-user@ip-10-0-0-71 ~]$ hadoop fs -ls /user/ec2-user/time_kpi/
Found 33 items
drwxr-xr-x - ec2-user ec2-user 0 2021-12-31 11:29 /user/ec2-user/time_kpi/_spark_metadata
-rw-r--r-- 3 ec2-user ec2-user 208 2021-12-31 11:22 /user/ec2-user/time_kpi/part-00000-01ac36a-f1d0-4654-87ad-381d013da283-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:17 /user/ec2-user/time_kpi/part-00000-2af4f700-b50a-4243-a0ef-013bb97f3337-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:26 /user/ec2-user/time_kpi/part-00000-32cd3877-f235-450a-b9bd-1cf1b76744f-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:24 /user/ec2-user/time_kpi/part-00000-64d5efe0-c675-4018-8228-5ba62e6d14bb-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:16 /user/ec2-user/time_kpi/part-00000-76a3b827-3225-4f83-9a8e-715754879af6-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:13 /user/ec2-user/time_kpi/part-00000-fa2e6e37-9e63-4b03-9099-377dd1bc3331-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:15 /user/ec2-user/time_kpi/part-00000-7dc12ee7-f1ff-4e11-0e04-d3903fe953e-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:25 /user/ec2-user/time_kpi/part-00000-f1b5f11-8014-413d-4467-1267ace19e9b-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:12 /user/ec2-user/time_kpi/part-00000-9502dbc0-c1dc-4d13-9a17-d9a54996094d-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:19 /user/ec2-user/time_kpi/part-00000-bd5be4ae-cdd4-4527-9cfb-bb14b18767c0-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:30 /user/ec2-user/time_kpi/part-00000-d7f33738-c99c-44cc-0265-102db3adffbb-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:28 /user/ec2-user/time_kpi/part-00000-e32777e2-fa85-404a-9c9e-f4a37709e32e-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:23 /user/ec2-user/time_kpi/part-00000-c3d2c943-f611-40d3-862a-000c608159e6-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:27 /user/ec2-user/time_kpi/part-00000-ea751707-aal2-473e-c051-752345dd2bc6-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:20 /user/ec2-user/time_kpi/part-00000-c6f143b17-1441-4112-9a28-2984b11bc910-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:29 /user/ec2-user/time_kpi/part-00000-f275ea00-75b6-4715-8eb8-95fa7b9ab0c4-c000.json
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:11 /user/ec2-user/time_kpi/part-00000-fc80b09c-e7df-40ff-9b48-dc4e8a09a5f2-c000.json
-rw-r--r-- 3 ec2-user ec2-user 211 2021-12-31 11:28 /user/ec2-user/time_kpi/part-00011-c0535c6a-cf82-4e93-82b8-9367c19e3c1c-c000.json
-rw-r--r-- 3 ec2-user ec2-user 194 2021-12-31 11:23 /user/ec2-user/time_kpi/part-00026-7475d6b5-877f-41dd-998c-ea20821b2d53-c000.json
-rw-r--r-- 3 ec2-user ec2-user 193 2021-12-31 11:20 /user/ec2-user/time_kpi/part-00056-b2d5bca2-c058-4183-81ca-b5d1c3c32f19-c000.json
-rw-r--r-- 3 ec2-user ec2-user 193 2021-12-31 11:17 /user/ec2-user/time_kpi/part-00056-faf34964-d719-4d22-a032-7613657c9a9a-c000.json
-rw-r--r-- 3 ec2-user ec2-user 211 2021-12-31 11:20 /user/ec2-user/time_kpi/part-00066-ac53a34d-63da-443a-906a-65acd2b6b340-c000.json
-rw-r--r-- 3 ec2-user ec2-user 194 2021-12-31 11:16 /user/ec2-user/time_kpi/part-00074-9dbd44d6-c5f3-4de8-adc7-1b8c36609c27-c000.json
-rw-r--r-- 3 ec2-user ec2-user 194 2021-12-31 11:23 /user/ec2-user/time_kpi/part-00107-edbd248b-9312-4b07-875b-3391bfedc57a-c000.json
-rw-r--r-- 3 ec2-user ec2-user 195 2021-12-31 11:22 /user/ec2-user/time_kpi/part-00123-bc90a047-ba20-47aa-a42e-5a9fc3b2c8c1-c000.json
-rw-r--r-- 3 ec2-user ec2-user 211 2021-12-31 11:16 /user/ec2-user/time_kpi/part-00145-24dbacd-dabd-49bf-bdfe-97e082cc077e-c000.json
-rw-r--r-- 3 ec2-user ec2-user 195 2021-12-31 11:22 /user/ec2-user/time_kpi/part-00150-fa473b87-37e0-4e6a-ac8b-3a05a2a2b1aa-c000.json
-rw-r--r-- 3 ec2-user ec2-user 211 2021-12-31 11:24 /user/ec2-user/time_kpi/part-00150-92e0322e-b5f0-44f3-9d8a-a7d57fbd956e-c000.json
-rw-r--r-- 3 ec2-user ec2-user 194 2021-12-31 11:22 /user/ec2-user/time_kpi/part-00160-c332372c-c961-4411-8d7a-1b1b049830d-c000.json
-rw-r--r-- 3 ec2-user ec2-user 210 2021-12-31 11:29 /user/ec2-user/time_kpi/part-00161-929c3158-e728-4139-a0c1-a0affb448d99-c000.json
-rw-r--r-- 3 ec2-user ec2-user 211 2021-12-31 11:27 /user/ec2-user/time_kpi/part-00171-3060d19e-a24d-434b-88a2-f423c9feaf67-c000.json
-rw-r--r-- 3 ec2-user ec2-user 196 2021-12-31 11:25 /user/ec2-user/time_kpi/part-00178-9bca5fd3-9567-44e8-9a79-35d4f053501f-c000.json
```

hadoop fs -ls /user/ec2-user/country\_kpi/

```
[ec2-user@ip-10-0-0-71 ~]$ hadoop fs -ls /user/ec2-user/country_kpi/
Found 47 items
drwxr-xr-x - ec2-user ec2-user 0 2021-12-31 11:29 /user/ec2-user/country_kpi/_spark_metadata
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:12 /user/ec2-user/country_kpi/part-00000-02ca18734-9980-44d1-915c-87318612c3b5-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:26 /user/ec2-user/country_kpi/part-00000-1d0mm61-8d17-4448-8af6-c68c9c8b4d6d-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:27 /user/ec2-user/country_kpi/part-00000-21a2mm6a-2923-458a-ba17-741370c10du2-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:21 /user/ec2-user/country_kpi/part-00000-303d00c6-7199-41e1-b06e-94cc8a28632a-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:20 /user/ec2-user/country_kpi/part-00000-31cc942d-7841-4b61-ae34-1b33cc827d92-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:28 /user/ec2-user/country_kpi/part-00000-4b46a1d-8377-41c8-9ee4-a06tab19dd6-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:11 /user/ec2-user/country_kpi/part-00000-4a958049-e1a6-4dd1-bced-c37dadd74878-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:13 /user/ec2-user/country_kpi/part-00000-62cc8ca9-d570-4da6-9732-c328e98f45b3-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:16 /user/ec2-user/country_kpi/part-00000-89afe87a-34b4-47db-a053-dbb0884e0afb-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:24 /user/ec2-user/country_kpi/part-00000-948cf9d0-4ea0-4476-8053-d5ad5f236bae-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:17 /user/ec2-user/country_kpi/part-00000-9ef38be4-c9c0-402d-8ce7-0582ec9e2a2b-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:22 /user/ec2-user/country_kpi/part-00000-a95ccabe-a71c-4b7d-0c37-072732e45b79-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:25 /user/ec2-user/country_kpi/part-00000-a90329fe-0ddc-4f1e-0d30-4fd047af2177-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:29 /user/ec2-user/country_kpi/part-00000-aed3525a-0f99-40da-b450-cc5cdc12cef0-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:30 /user/ec2-user/country_kpi/part-00000-bbed1f42-9644-41b3-a25d-b9fcaa39a953-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:14 /user/ec2-user/country_kpi/part-00000-d50fe2a6-5c22-47f2-a667-75df4b233a4c-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 0 2021-12-31 11:23 /user/ec2-user/country_kpi/part-00000-da302d2f-bc09-4020-adb9-363b0906533f-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 173 2021-12-31 11:23 /user/ec2-user/country_kpi/part-00001-fe3490d9-74cc-444c-b114-120c24c11214-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 165 2021-12-31 11:30 /user/ec2-user/country_kpi/part-00002-0909566f-9375-4e21-b8d0-5fc14c128968-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 169 2021-12-31 11:28 /user/ec2-user/country_kpi/part-00008-d5e71841-abee-4003-8bd4-000c61d5b846-c000.json
OR
-rw-r--r-- 3 ec2-user ec2-user 175 2021-12-31 11:24 /user/ec2-user/country_kpi/part-00016-772a7ea3-7b15-4c64-a5b8-40ca31e7ab98-c000.json
```

And used 'cat' command to take a look at the data

hadoop fs -cat /user/ec2-user/time\_kpi/part\*



```
[ec2-user@ip-10-0-0-71 ~]$ hadoop fs -cat /user/ec2-user/time_kpi/part*
{"window":{"start":"2021-10-25T11:10:00.000Z","end":"2021-12-31T11:16:00.000Z"},"OPM":9,"total_sale_volume":245.0800018310547,"average_transaction_size":27.23111131456163,"rate_of_return":0.1111111111111111}
{"window":{"start":"2021-10-25T11:24:00.000Z","end":"2021-12-31T11:25:00.000Z"},"OPM":12,"total_sale_volume":1079.4999951422215,"average_transaction_size":99.95033292951046,"rate_of_return":0.003333333333333333}
{"window":{"start":"2021-10-25T11:19:00.000Z","end":"2021-12-31T11:20:00.000Z"},"OPM":11,"total_sale_volume":428.7600100636482,"average_transaction_size":39.97010273305933,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:13:00.000Z","end":"2021-12-31T11:14:00.000Z"},"OPM":9,"total_sale_volume":459.8600026177002,"average_transaction_size":51.09555503530002,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:13:00.000Z","end":"2021-12-31T11:14:00.000Z"},"OPM":9,"total_sale_volume":459.8600026177002,"average_transaction_size":51.09555503530002,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:14:00.000Z","end":"2021-12-31T11:15:00.000Z"},"OPM":14,"total_sale_volume":483.4499945640564,"average_transaction_size":34.532142408961174,"rate_of_return":0.07142857142857142}
{"window":{"start":"2021-10-25T11:11:00.000Z","end":"2021-12-31T11:12:00.000Z"},"OPM":5,"total_sale_volume":396.11000061035156,"average_transaction_size":79.22200012207031,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:20:00.000Z","end":"2021-12-31T11:21:00.000Z"},"OPM":5,"total_sale_volume":250.5600032063965,"average_transaction_size":51.7200066612794,"rate_of_return":0.22}
{"window":{"start":"2021-10-25T11:17:00.000Z","end":"2021-12-31T11:18:00.000Z"},"OPM":7,"total_sale_volume":1067.4699907302056,"average_transaction_size":52.4951796145937,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:12:00.000Z","end":"2021-12-31T11:13:00.000Z"},"OPM":13,"total_sale_volume":460.60000121593475,"average_transaction_size":36.04615393968729,"rate_of_return":0.07692307692307693}
{"window":{"start":"2021-10-25T11:18:00.000Z","end":"2021-12-31T11:19:00.000Z"},"OPM":7,"total_sale_volume":1583.7199920415878,"average_transaction_size":226.24571314379827,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:21:00.000Z","end":"2021-12-31T11:22:00.000Z"},"OPM":14,"total_sale_volume":668.9900054335594,"average_transaction_size":47.78500039811385,"rate_of_return":0.14285714285714285}
{"window":{"start":"2021-10-25T11:16:00.000Z","end":"2021-12-31T11:17:00.000Z"},"OPM":4,"total_sale_volume":78.17999869385037,"average_transaction_size":19.544999673962593,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:25:00.000Z","end":"2021-12-31T11:26:00.000Z"},"OPM":15,"total_sale_volume":542.6700085401535,"average_transaction_size":36.17000056934357,"rate_of_return":0.13333333333333333}
{"window":{"start":"2021-10-25T11:23:00.000Z","end":"2021-12-31T11:24:00.000Z"},"OPM":15,"total_sale_volume":469.79999724030495,"average_transaction_size":31.31999901502033,"rate_of_return":0.06666666666666667}
{"window":{"start":"2021-10-25T11:22:00.000Z","end":"2021-12-31T11:23:00.000Z"},"OPM":12,"total_sale_volume":335.83000135421753,"average_transaction_size":27.995033446104795,"rate_of_return":0.0}
```

hadoop fs -cat /user/ec2-user/country\_kpi/part\*

```
[ec2-user@ip-10-0-0-71 ~]$ hadoop fs -cat /user/ec2-user/country_kpi/part*
{"window":{"start":"2021-10-25T11:20:00.000Z","end":"2021-12-31T11:21:00.000Z"},"country":"Unspecified","OPM":1,"total_sale_volume":93.43000030517578,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:26:00.000Z","end":"2021-12-31T11:27:00.000Z"},"country":"EIRE","OPM":1,"total_sale_volume":29.15999904741211,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:24:00.000Z","end":"2021-12-31T11:25:00.000Z"},"country":"Denmark","OPM":1,"total_sale_volume":157.77999877929688,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:21:00.000Z","end":"2021-12-31T11:22:00.000Z"},"country":"Channel Islands","OPM":1,"total_sale_volume":23.1299991607666,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:19:00.000Z","end":"2021-12-31T11:20:00.000Z"},"country":"United Kingdom","OPM":11,"total_sale_volume":428.7600100636482,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:23:00.000Z","end":"2021-12-31T11:24:00.000Z"},"country":"United Kingdom","OPM":15,"total_sale_volume":469.79999724030495,"rate_of_return":0.06666666666666667}
{"window":{"start":"2021-10-25T11:25:00.000Z","end":"2021-12-31T11:26:00.000Z"},"country":"France","OPM":2,"total_sale_volume":30.779999136924744,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:23:00.000Z","end":"2021-12-31T11:23:00.000Z"},"country":"Spain","OPM":1,"total_sale_volume":47.06999969402422,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:26:00.000Z","end":"2021-12-31T11:27:00.000Z"},"country":"United Kingdom","OPM":4,"total_sale_volume":228.5200045776367,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:17:00.000Z","end":"2021-12-31T11:18:00.000Z"},"country":"Norway","OPM":1,"total_sale_volume":11.899999618530273,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:12:00.000Z","end":"2021-12-31T11:13:00.000Z"},"country":"ETRE","OPM":1,"total_sale_volume":4.440000057220459,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:14:00.000Z","end":"2021-12-31T11:15:00.000Z"},"country":"United Kingdom","OPM":14,"total_sale_volume":483.4499945640564,"rate_of_return":0.07142857142857142}
{"window":{"start":"2021-10-25T11:13:00.000Z","end":"2021-12-31T11:14:00.000Z"},"country":"Germany","OPM":1,"total_sale_volume":30.59000015250789,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:12:00.000Z","end":"2021-12-31T11:13:00.000Z"},"country":"United Kingdom","OPM":12,"total_sale_volume":464.160011587143,"rate_of_return":0.06333333333333333}
{"window":{"start":"2021-10-25T11:21:00.000Z","end":"2021-12-31T11:22:00.000Z"},"country":"United Kingdom","OPM":13,"total_sale_volume":645.860062727920,"rate_of_return":0.15304615304615305}
{"window":{"start":"2021-10-25T11:18:00.000Z","end":"2021-12-31T11:19:00.000Z"},"country":"United Kingdom","OPM":7,"total_sale_volume":1583.7199920415878,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:22:00.000Z","end":"2021-12-31T11:23:00.000Z"},"country":"United Kingdom","OPM":11,"total_sale_volume":288.760016593933,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:13:00.000Z","end":"2021-12-31T11:14:00.000Z"},"country":"United Kingdom","OPM":7,"total_sale_volume":349.870008399234,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:16:00.000Z","end":"2021-12-31T11:17:00.000Z"},"country":"United Kingdom","OPM":4,"total_sale_volume":78.17999869385037,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:17:00.000Z","end":"2021-12-31T11:18:00.000Z"},"country":"United Kingdom","OPM":6,"total_sale_volume":1055.569901117554,"rate_of_return":0.0}
{"window":{"start":"2021-10-25T11:24:00.000Z","end":"2021-12-31T11:25:00.000Z"},"country":"United Kingdom","OPM":10,"total_sale_volume":930.2199963629246,"rate_of_return":0.0}
```

## Transfer of files from CDH Instance on AWS to my system, using WinSCP

First, I needed to transfer the JSON files from HDFS into the the EC2 system

I created directories for time-based and then time-and-country-based KPIs as ec2-user. Using the 'get' command I copied the contents of the output folders into the EC2 system.

```
mkdir timebased-KPI
```

```
hadoop fs -get /user/ec2-user/time_kpi /home/ec2-user/timebased-KPI
```

```
mkdir country-and-timebased-KPI  
hadoop fs -get /user/ec2-user/country_kpi /home/ec2-user/country-and-timebased-KPI
```

Thereafter I used WinSCP to establish a connection between the EC2 instance and my local file system to transfer all the required files into my system.